

# Use Case Tutorial

Gary Chen  
2018/10/23



CloudMile

# Agenda

Overall Workflow

Exploratory Data Analysis

Feature Engineering + Training

# Agenda

Overall Workflow

Exploratory Data Analysis

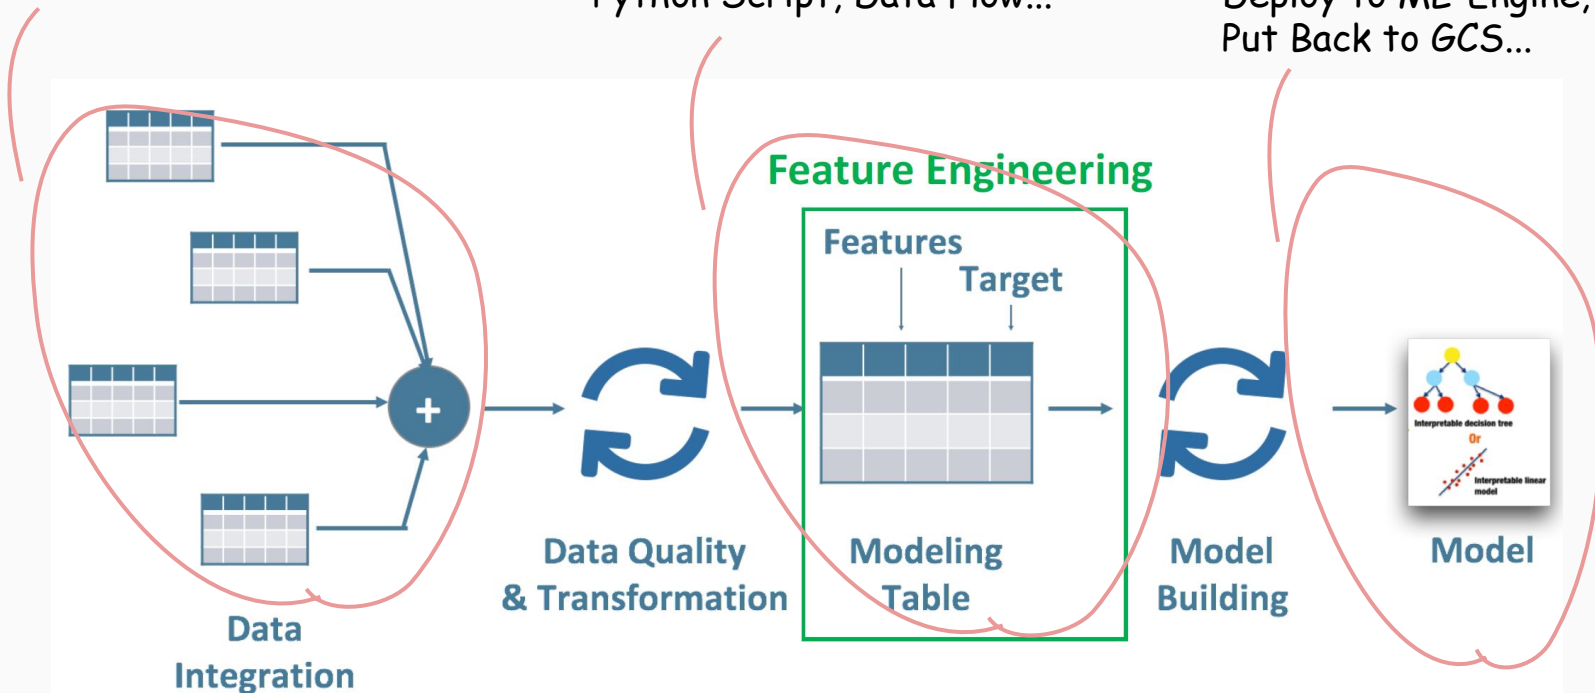
Feature Engineering + Training

# Typical Enterprise Machine Learning Workflow

CSV(GCS), Cloud SQL, Big Query

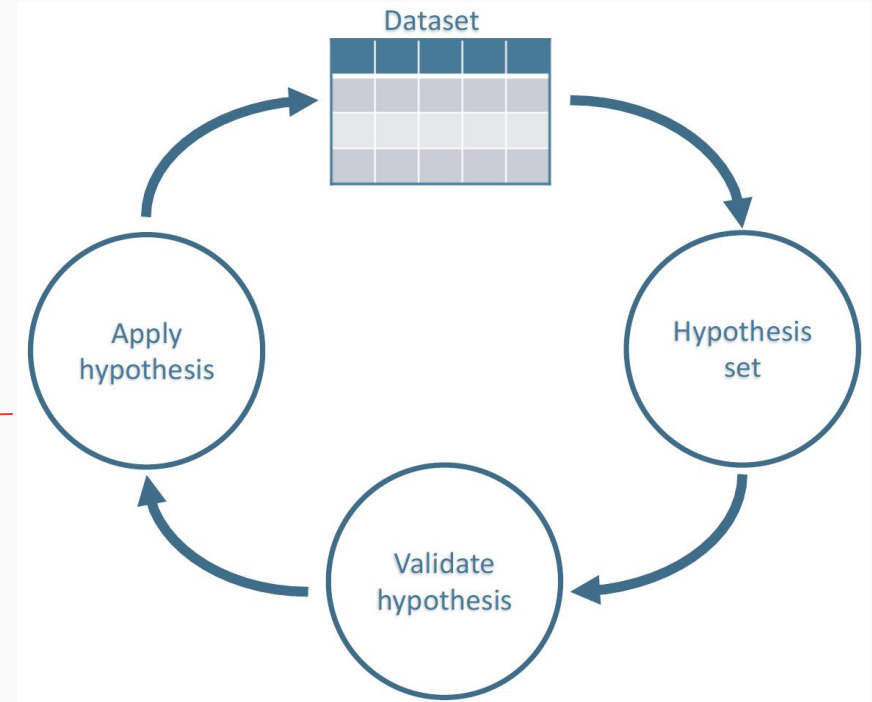
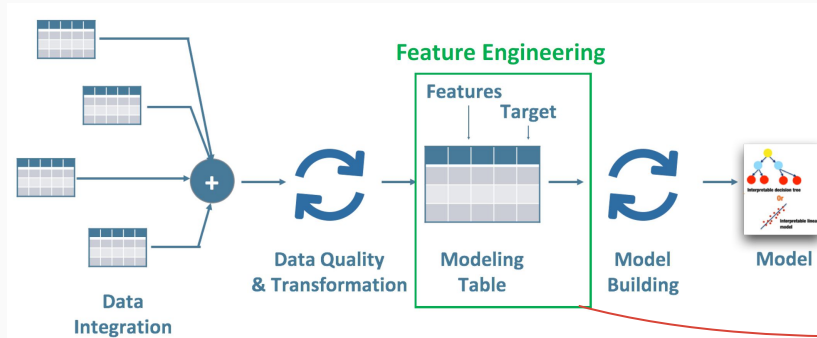
Python Script, Data Flow...

Deploy to ML-Engine, Put Back to GCS...



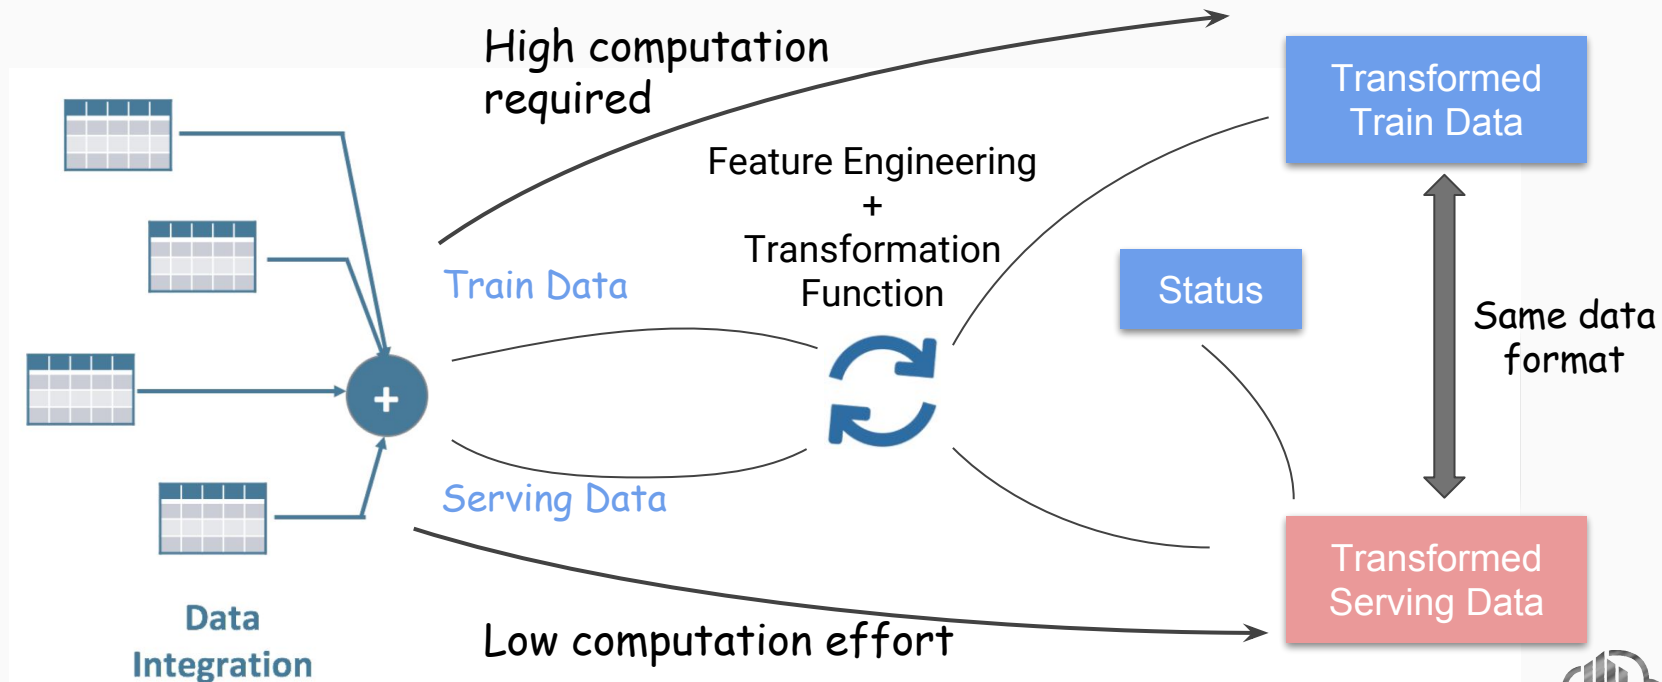
# Typical Enterprise Machine Learning Workflow

## Feature Engineering cycle



# Typical Enterprise Machine Learning Workflow

## The Training and Serving



# Agenda

Overall Workflow

Exploratory Data Analysis

Feature Engineering + Training

# Type of Variable

## ❑ Numerical variable: “Float, Integer”

- ❑ Discrete numerical variable

e.g: [1, 2, 3, 4]

- ❑ Continuous numerical variable

e.g: [1.23, 0.87, 1.5498, -0.3146]

## ❑ Nominal (Categorical) variable: “String, Integer”

e.g: Geography: ['France', 'Germany', 'Spain']

e.g: Address:

1F., No.1, Bilong Ln., Zhongzheng 1st Rd., Yingge Dist., New Taipei City 239, Taiwan (R.O.C.)


## ❑ Ordinal nominal variable: “String, Integer”

e.g: Size of clothes: ['S', 'M', 'L', 'XL']



# What We Want to Explore

- ❑ Numerical variables:  
mean, std, median, quartiles, deciles  
data distribution (histogram)
- ❑ Nominal variables: frequency distribution
- ❑ Relationship between variables
  - ❑ Numerical x Numerical
  - ❑ Numerical x Nominal
  - ❑ Nominal x Nominal




What we care is the relation between **Feature** and **Label**

# EDA Tools


## pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$




Data read write, data clean, data transformation, data join ...

## seaborn: statistical data visualization



Data visualization



Statistical test, basic model



Numerical Variable Visualized

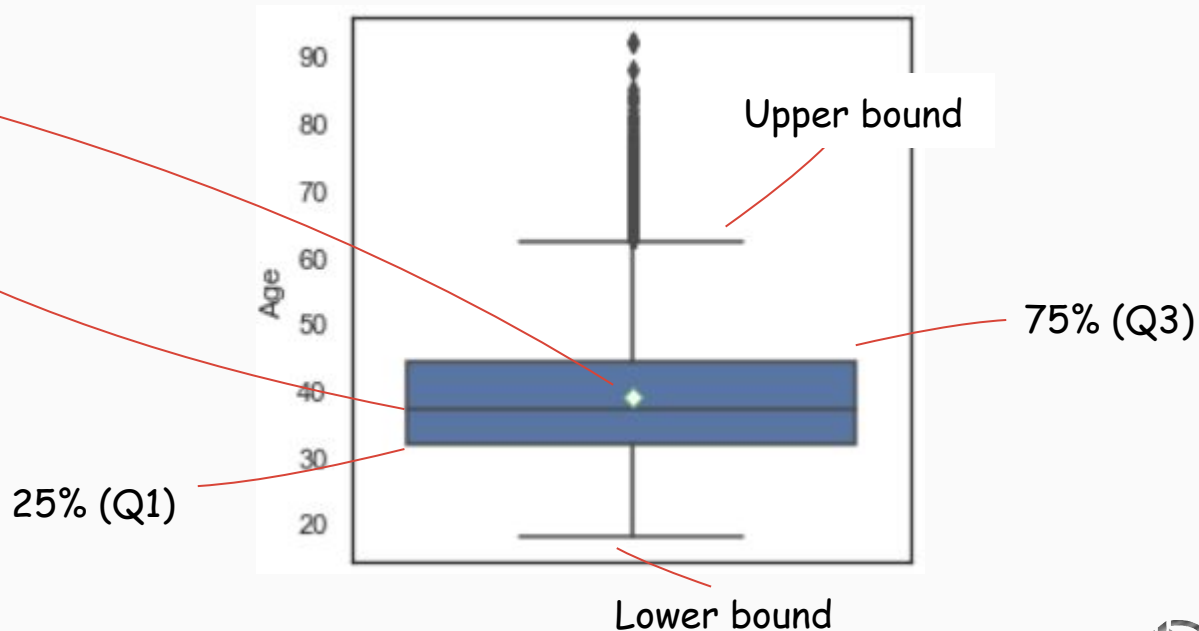
# Visualization - Numerical Variable (Boxplot)

```
pandas.Series.describe()
```

```
raw.Age.describe()
```

```
count    7500.000000  
mean      39.004267  
std       10.500007  
min       18.000000  
25%       32.000000  
50%       37.000000  
75%       44.000000  
max       92.000000  
Name: Age, dtype: float64
```

```
sns.boxplot(raw.Age,  
            showmeans=True, orient='v',  
            meanprops={'marker': 'D', 'markerfacecolor': 'white'})
```




# Visualization - Numerical Variable (Boxplot)

Interquartile range

$$\text{IQR} = Q3 - Q1$$

$$\text{Upper bound} = Q3 + \text{IQR} \times 1.5$$

$$\text{Lower bound} = Q1 - \text{IQR} \times 1.5$$




# Visualization - Numerical Variable (Violinplot)

Include information


- Boxplot
- Data distribution

```
sns.violinplot(raw.Age, orient='v')
```




# Visualization - Numerical Variable (Distplot, Lineplot)

```
sns.distplot(raw.Age)
```



```
sns.lineplot(  
    np.arange(len(raw)),  
    raw.Age.sort_values())
```






Nominal Variable Visualized



# Visualization - Nominal Variable (Countplot)

```
sns.countplot(  
    x="AgeBin",  
    data=raw)
```



```
pandas.Series.value_counts
```

```
raw.AgeBin.value_counts()
```

14-30	1209.0
30-35	1525.0
35-40	1721.0
40-45	1267.0
45-62	1465.0
62up	313.0
Name: AgeBin, dtype: float64	




## Multivariate Visualized

# Multivariate Visualized - Numerical x Numerical

## Regression Plot

⇒ sns.lmplot


```
g = sns.lmplot(x='Age', y='CreditScore',  
               # hue='Exited',  
               data=raw, height=4, aspect=1.5, markers='x')
```



# Multivariate Visualized - Nominal x Nominal

Heatmap  $\Rightarrow$  sns.heatmap

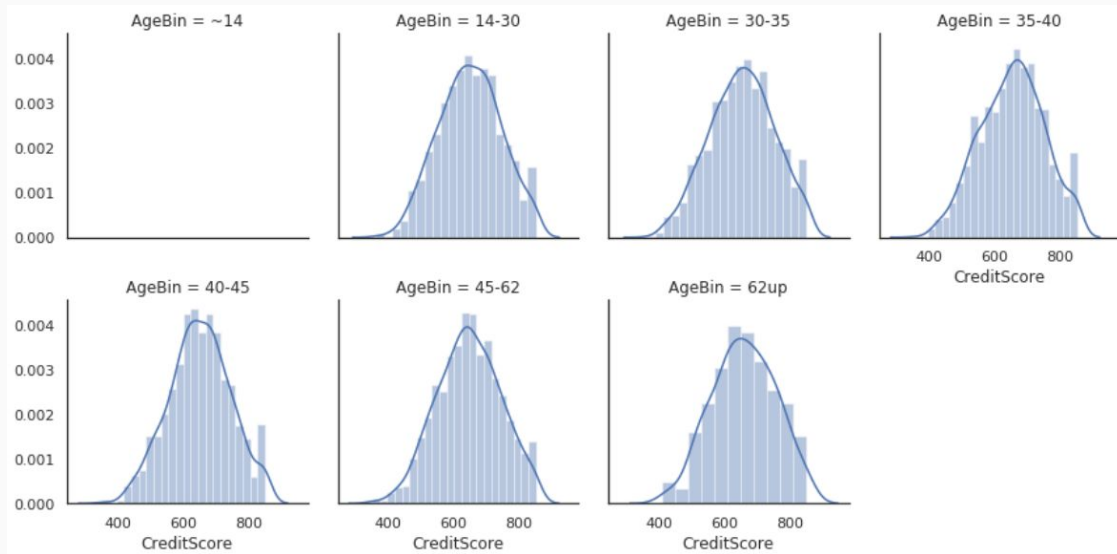

```
heatmap(raw, 'Geography', 'Gender', target='Exited')
```



# Multivariate Visualized - Nominal x Numerical

Divide numerical variable by nominal variable !

⇒ Violinplot, Boxplot, Distplot, Lineplot ...





## About The Statistical

# Numerical x Numerical (Linear)

## Pearson Correlation Coefficient

$$\rho = \frac{\text{cov}(X, Y)}{\sigma_x \sigma_y} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}} \quad -1 < \rho < 1$$


```
pd.DataFrame.corr()
```

```
data = data[  
    ['Tenure', 'Balance', 'EstimatedSalary']  
]  
data.corr(method='pearson')
```

	Tenure	Balance	EstimatedSalary
Tenure	1.000000	-0.012194	0.014443
Balance	-0.012194	1.000000	0.010461
EstimatedSalary	0.014443	0.010461	1.000000

# Numerical x Numerical (Linear)

## Pearson Correlation Coefficient





# Numerical x Numerical (Linear)

## Spearman's rank correlation coefficient

$$\rho = 1 - \frac{6 \sum d_i^2}{n(n^2 - 1)}$$

$$-1 < \rho < 1$$

Difference between  
rank of two variables

```
pd.DataFrame.corr()
```

```
data = data[  
    ['Tenure', 'Balance', 'EstimatedSalary']  
]  
data.corr(method='spearman')
```

Data size

	Tenure	Balance	EstimatedSalary
Tenure	1.000000	-0.008285	0.014451
Balance	-0.008285	1.000000	0.006840
EstimatedSalary	0.014451	0.006840	1.000000

# Numerical x Numerical (Linear)

## Spearman's rank correlation coefficient

	x	y	x_level	y_level	d_i	d_i^2
0	86	0	1	1	0	0
1	97	20	2	6	-4	16
2	99	28	3	8	-5	25
3	100	27	4	7	-3	9
4	101	50	5	10	-5	25
5	103	29	6	9	-3	9
6	106	7	7	3	4	16
7	110	17	8	5	3	9
8	112	6	9	2	7	49
9	113	12	10	4	6	36

$$\rho = 1 - \frac{6 \times 194}{10(10^2 - 1)} = -0.175175$$

Negative correlation,  
but not significant

Beware the **scaler information lost**,  
good for ordinal variable



## Chi-Square Test


# About The Statistical Test

## Before The Chi-Square Test ..., About The Statistical Experiment

1. Hypothesis
  - a.  $H_0$ : Null Hypothesis  
e.g: Variable  $X_1$ ,  $X_2$  independent
  - b.  $H_1$ : Alternative Hypothesis  
e.g: Variable  $X_1$  association with  $X_2$
2. Significant Level  $\alpha$ , Confidence Interval  $(1 - \alpha)$   
e.g: Given  $\alpha = 0.05$ , means confidence interval = **0.95**
3. The p-value
4. When to reject  $H_0$ ? (Means the result is significant)  
 $H_0$ : Negative event  
 $H_1$ : Positive event



# About The Statistical Test




		Prediction	
		$H_0$	$H_1$
Ground Truth	$H_0$	TN	FP
	$H_1$	FN	TP

Type I error  $\Rightarrow$  FP  
Type II error  $\Rightarrow$  FN

$\alpha$ : Significant level  
 $\Rightarrow$  The probability of type I error occurred  
[Example of Type I II error](#)

# About The Statistical Test

When is the experiment significant?



**PDF:** Probability Density Function

**$t_0$ :** Statistic value

**p-value:** Area under **PDF** and  **$t_0$**


$P < \alpha$ : Reject  $H_0$

$P \geq \alpha$ : Not reject  $H_0$

# About The Statistical Test

When is the experiment significant?

PDF:  $y = f(x)$



**PDF:** Probability Density Function

**$t_0$ :** Statistic value

$t_0 > t'$ : Reject  $H_0$

$t_0 \leq t'$ : Not reject  $H_0$

# About The Statistical Test

- Goodness of Fit
  - ◆  $H_0$ : There is no difference between the observed and expected frequencies
  - ◆  $H_1$ : There is a difference between the observed and the expected frequencies
  
- Test for homogeneity
  - ◆  $H_0$ : Populations follow the same probability distribution
  - ◆  $H_1$ : One of populations doesn't follow the specific probability distribution
  
- Test for Independent
  - ◆  $H_0$ : Nominal variables  $X_1, X_2$  independent
  - ◆  $H_1$ : Nominal variables  $X_1$  association with  $X_2$



# About The Statistical Test

H<sub>0</sub>: Variable **X<sub>1</sub>**, **X<sub>2</sub>** independent

H<sub>1</sub>: Variable **X<sub>1</sub>** association with **X<sub>2</sub>**


marit	educ
Never married	PhD or higher
Married	Middle school or lower
Divorced	Bachelor's
Widowed	PhD or higher
Married	PhD or higher

Marital Status by Education | n = 300

	Middle school or lower	High school	Bachelor's	Master's	PhD or higher	Total
Never married	18	36	21	9	6	90
Married	12	36	45	36	21	150
Divorced	6	9	9	3	3	30
Widowed	3	9	9	6	3	30
Total	39	90	84	54	33	300

$$\chi^2 = \sum_i \frac{(Actual_i - Expected_i)^2}{Expected_i}$$

t<sub>0</sub>: Chi-Square  
Statistic Value



# Nominal x Nominal (Non-Linear): Chi-Square Test

H<sub>0</sub>: Variable **X<sub>1</sub>**, **X<sub>2</sub>** independent    H<sub>1</sub>: Variable **X<sub>1</sub>** association with **X<sub>2</sub>**

Marital Status by Education | n = 300

	Middle school or lower	High school	Bachelor's	Master's	PhD or higher	Total
Never married	18	36	21	9	6	90
Married	12	36	45	36	21	150
Divorced	6	9	9	3	3	30
Widowed	3	9	9	6	3	30
Total	39	90	84	54	33	300

$$\chi^2 = \sum_i \frac{(Actual_i - Expected_i)^2}{Expected_i}$$

Assume the events between **Education** and **Marital** status are mutual independent

$$P(A \cap B) = P(A)P(B)$$

$$P(\mathbf{Education} \text{ and } \mathbf{Marital}) = P(\mathbf{Education}) \times P(\mathbf{Marital})$$

# Nominal x Nominal (Non-Linear): Chi-Square Test

$H_0$ : Variable  $X_1$ ,  $X_2$  independent     $H_1$ : Variable  $X_1$  association with  $X_2$

Marital Status by Education | n = 300

	Middle school or lower	High school	Bachelor's	Master's	PhD or higher	Total
Never married	18	36	21	9	6	90
Married	12	36	45	36	21	150
Divorced	6	9	9	3	3	30
Widowed	3	9	9	6	3	30
Total	39	90	84	54	33	300

$$\chi^2 = \sum_i \frac{(Actual_i - Expected_i)^2}{Expected_i}$$

Actual(**Master's** x **Married**) = 36

Expected(Master's x Married) =  $P(\text{Master's}) \times P(\text{Married}) \times \text{Length}(\text{data})$   
=  $(54 / 300) \times (150 / 300) \times 300$   
=  $0.18 \times 0.5 \times 300$   
= 27

$\Rightarrow (36 - 27)^2 / 27 = 3.0$

# Nominal x Nominal (Non-Linear): Chi-Square Test

## statsmodels package example

```
# Calculate chi-square value, p-value, degree of freedom, expected value  
chi, pv, df, expected = stats.chi2_contingency(observed=data)  
  
# check if chi-square value > criterion(95% confidence interval)  
crit = stats.chi2.ppf(q=0.95, df=df)  
  
print(f'chi-square value: {chi}, criterion: {crit}')
```

chi-square value: 1022.025, criterion: 11.070  
result: True

Significant! Reject  $H_0$   
There is a relationship between  $X_1$ ,  $X_2$



## ANOVA (Analysis of Variance)



# Nominal x Numerical (Non-Linear): ANOVA


Assume we have k group,  $k > 1$

$H_0: \mu_1 = \mu_2 \dots = \mu_k$

$H_1$ : Means are not all equal.

Prerequisite

- Each groups approximately follow **normal distribution (Guassian distribution)**
- Independent cases
- Equality (or "homogeneity") of variances



Welch test, Brown Forsythe test...

??% Confidence in ANOVA test

# Nominal x Numerical (Non-Linear): ANOVA

Assume we have k group,  $k > 1$

$H_0$ :  $\mu_1 = \mu_2 \dots = \mu_k$


$H_1$ : Means are not all equal.

Prerequisite


- Each groups approximately follow **normal distribution (Guassian distribution)**
- Independent cases
- Equality (or "homogeneity") of variances

**X**: Numerical variable

**Y** : Nominal variable



# Nominal x Numerical (Non-Linear): ANOVA



Total sum of square

$$SS_T = \sum_i (X_i - \overline{X_{all}})^2$$

Sum of square between groups

$$SS_B = \sum_{\#group} (\overline{X_{group}} - \overline{X_{all}})^2$$

Sum of square within groups

$$SS_W = \sum_{\#group} \sum_{\#within\_group} (X_i - \overline{X_{group}})^2$$

$$SST = SSB + SSW$$



# Nominal x Numerical (Non-Linear): ANOVA

$$MS_B = \frac{SS_B}{df_B} \quad \rightarrow \quad \frac{MS_B}{MS_W}$$

$$MS_W = \frac{SS_W}{df_W}$$

	SS	DF	MS	F	P
<b>B</b>	SS <sub>B</sub>	G - 1	MS <sub>B</sub>	MS <sub>B</sub> / MS <sub>W</sub>	P-value
<b>W</b>	SS <sub>W</sub>	(N - 1) - (G - 1)	MS <sub>W</sub>		
<b>T</b>	SS <sub>T</sub>	(N - 1)			

We hope the F Larger

Area =>  $P < \alpha$ : Reject  $H_0$   
 $P \geq \alpha$ : Not reject  $H_0$

Total sum of square

$$SS_T = \sum_i (X_i - \overline{X_{all}})^2$$

Sum of square between groups

$$SS_B = \sum_{\#group} (\overline{X_{group}} - \overline{X_{all}})^2$$

Sum of square within groups

$$SS_W = \sum_{\#group} \sum_{\#within\_group} (X_i - \overline{X_{group}})^2$$



# Nominal x Numerical (Non-Linear): ANOVA

## statsmodels package example

```
def anova(formula, data):  
    table = sm.stats.anova_lm(  
        ols(formula, data=data).fit(),  
        typ=2  
    )  
    return table  
  
anova("Age ~ C(Exited)", data=data)
```

"Numerical ~ C(Nominal)"

	sum_sq	df	F	PR(>F)
C(Exited)	63991.391600	1.0	629.02925	2.266897e-133
Residual	762774.471867	7498.0	NaN	NaN

P-value

# Recap

- ❑ Single Variable
  - ❑ Numerical variable  $\Rightarrow$  Boxplot, Violinplot, Distplot, Lineplot
  - ❑ Nominal variable  $\Rightarrow$  Countplot
- ❑ Multi-Variable
  - ❑ Numerical x Numerical  $\Rightarrow$  Regression plot
  - ❑ Nominal x Nominal  $\Rightarrow$  Heatmap
  - ❑ Nominal x Numerical  $\Rightarrow$  Conditional Boxplot, Conditional Violinplot ...
- ❑ Relationship
  - ❑ Numerical x Numerical  $\Rightarrow$  Pearson, Spearman Correlation Coefficient
  - ❑ Nominal x Nominal  $\Rightarrow$  Chi-Square
  - ❑ Nominal x Numerical  $\Rightarrow$  ANOVA

# Lab: Exploratory Data Analysis

## Topic : Exploratory Data Analysis

<b>Filename</b>	lab_eda_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	<ul style="list-style-type: none"><li>→ Understand the data</li><li>→ Features distribution</li><li>→ Relationship between features or between features and label</li></ul>
<b>Duration</b>	About 20 min

# Agenda


Overall Workflow

Exploratory Data Analysis

Feature Engineering + Training

# Basic - Data Clean

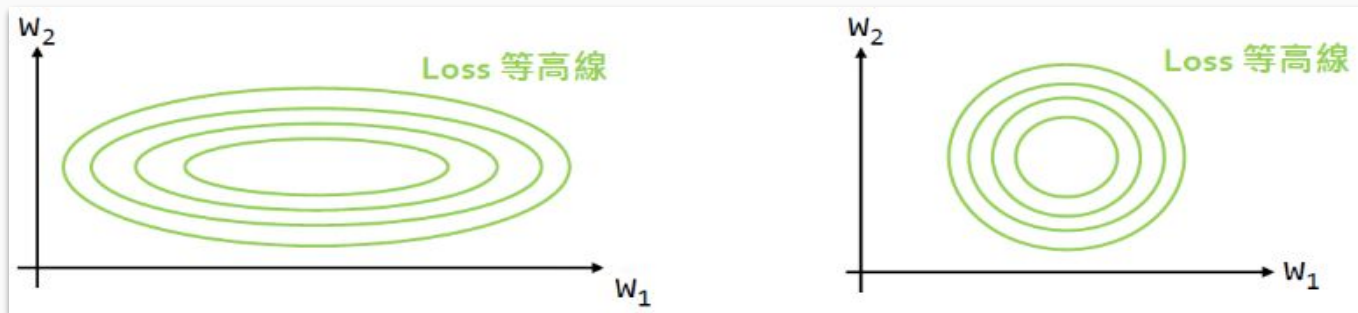

- Missing value in numerical variable
  - ◆ Fill **mean** or **median**
- Outlier in numerical variable
  - ◆ Utilize the quartile to find **outlier**, and fill mean or median
- Missing value in nominal variable
  - ◆ See missing value as an special class
  - ◆ Add a feature to describe missing value



nominal column	added
A	0
NaN	1
B	0

# Basic

Numeric variable  $\Rightarrow$  Normalization



# Basic

Numeric variable  $\Rightarrow$  Normalization

- Min max scaler to [0, 1] **(Beware outlier)**

$$\frac{x_i - \min}{\max - \min}$$


- Scale to standard normal distribution (Z-score standardize)

$$\mu = 0, \sigma = 1$$



# Basic

Nominal variable  $\Rightarrow$  One Hot Encoding



1
2
3



all the pairwise distance  
are the same  $\Rightarrow \sqrt{2}$



## Before Advanced Feature Engineering

# Before Advanced Feature Engineering

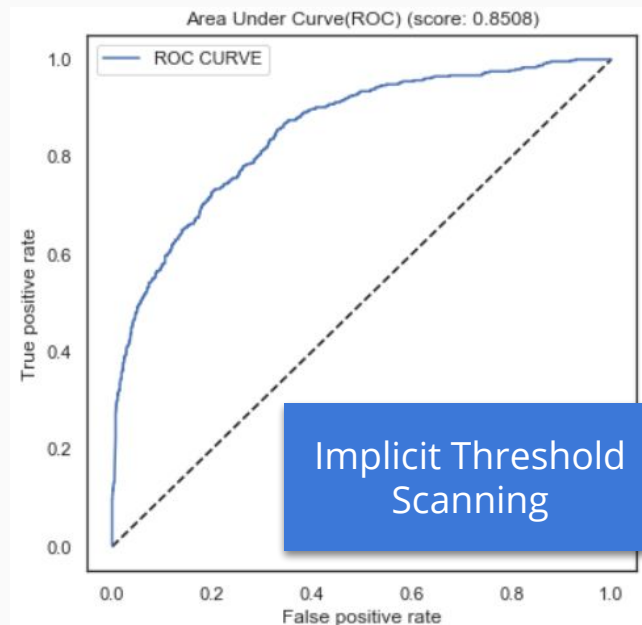
## Topic : Knowing Progaming Structure

<b>Filename</b>	lab_model_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	
<b>Duration</b>	

# Before Advanced Feature Engineering

## ROC: Receiver Operating Characteristic

TP Rate cross FP Rate  
( $0 < AUC < 1$ )



AUC = 0.5 (no discrimination)

$0.7 \leq AUC \leq 0.8$  (acceptable discrimination)

$0.8 \leq AUC \leq 0.9$  (excellent discrimination)

$0.9 \leq AUC \leq 1.0$  (outstanding discrimination)

		Prediction		
		0	1	
Ground Truth	0	TN	FP	FP Rate = $FP / (TN + FP)$
	1	FN	TP	

TP Rate =  $TP / (FN + TP)$

# Before Advanced Feature Engineering

## AUROC (Code)

```
from sklearn.metrics import roc_curve, auc  
fpr, tpr, thres = roc_curve(y, pred, pos_label=1)  
auc_scr = auc(fpr, tpr)
```

	fpr	tpr	threshold
0	0.000000	0.002045	0.999386
1	0.000000	0.104294	0.883024
2	0.000497	0.104294	0.880933
3	0.000497	0.110429	0.870801
4	0.000995	0.110429	0.870346
5	0.000995	0.118609	0.852421

# Before Advanced Feature Engineering

Find Best Threshold  $\Rightarrow$  F Beta Score

$$F_{\beta} = (1 + \beta^2) \cdot \frac{\text{precision} \cdot \text{recall}}{(\beta^2 \cdot \text{precision}) + \text{recall}}$$

Small  $\beta \Rightarrow$  prefer precision  
Large  $\beta \Rightarrow$  prefer recall

Suggested adjust range  
 $0.1 < \beta < 2$

		Prediction		
		0	1	
Ground Truth	0	TN	FP	precision
	1	FN	TP	
		recall		

	precision	recall	f1-score	support
0	0.88	0.96	0.92	2011
1	0.73	0.45	0.56	489
avg / total	0.85	0.86	0.85	2500

# Before Advanced Feature Engineering

## Function `f_beta_scann`

```
def f_beta_scann(y_true, y_pred, beta=0.5):  
    """F beta score 掃描找出最佳 threshold"""  
  
    y_pred = pd.Series(y_pred.ravel())  
    # 切割100等分, 尋找最佳 f beta score  
    bins = np.linspace(y_pred.min(), y_pred.max(), 100)  
    # 找出F beta score最高的點  
    result =  
        np.array([  
            precision_recall_fscore_support(  
                y_true=y_true,  
                y_pred=y_pred > thres, beta=beta)[2][1]  
            for thres in bins])  
  
    best_idx = result.argmax()  
    return bins[best_idx], result[best_idx]
```

	0	1
precision		
Recall		
F score		
???		



## Advanced Feature Engineering




Coming up with features is difficult, time consuming,  
requires expert knowledge.

**"Applied machine learning"** is basically feature  
engineering.

-Andrew Ng

# Advanced - Binning Numerical Variable

Binning  $\Rightarrow$  Find The Non-Linear Relationship



$$Y = WX + b$$



$$Y = W_1X_1 + W_2X_2 + W_3X_3 + b$$


$$Y = W_1X_1 + \quad + b$$

$$Y = \quad W_2X_2 \quad + b$$

$$Y = \quad W_3X_3 + b$$

# Advanced - Binning Numerical Variable

Binning Example  $\Rightarrow$  Quartile Cut



# Advanced - Binning Numerical Variable

## Binning Example $\Rightarrow$ Quartile Cut

```
def quartile_binning(x):  
    # Quartile cut  
    bins = np.percentile(x, range(0, 100, 25))[1:].tolist()  
    # IQR  
    iqr_x_150 = (bins[-1] - bins[0]) * 1.5  
  
    bins = [  
        bins[0] - iqr_x_150 + bins + [  
            bins[-1] + iqr_x_150  
        ]  
    ]  
    result = pd.Series(np digitize(x, bins)) \  
        .map(pd.Series([0, 1, 2, 3, 4, 0])).values  
    return result, bins
```

## Topic : Try Binning

<b>Filename</b>	lab_eda_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	Add binning features: → Age, HasBalance, CreditScore, Tenure, EstimatedSalary
<b>Duration</b>	About 10 min

# Advanced - WOE Encoding (Only for Binary Classification)

Weight of Evidence

Inspired from **Logistic Regression**  $\Rightarrow$  Binary Classification  $\Rightarrow$  0 or 1

$$\underline{\sigma(wx + b)} \quad z = wx + b$$

$$\sigma = \frac{1}{1 + e^{-z}}$$

**LR function**

$$\text{Odds Ratio} \Rightarrow \frac{P}{1 - P}$$

$$\frac{P}{1 - P} = \frac{\frac{1}{1 + e^{-z}}}{1 - \frac{1}{1 + e^{-z}}} = \dots = e^z = e^{wx + b}$$

$$\Rightarrow \log\left(\frac{P}{1 - P}\right) = wx + b$$

For the interpretable

# Advanced - WOE Encoding (Only for Binary Classification)

**Cancer Prediction Example**  $\Rightarrow wx + b = 0.095 \cdot age + 2.645$

Log Odds Ratio  $\log\left(\frac{P}{1-P}\right) = 0.095 \cdot age + 2.645$

Odds Ratio  $e^{\log\left(\frac{P}{1-P}\right)} = e^{0.095age+2.645} = e^{0.095age} e^{2.645}$

**If age increase 1**

$$e^{0.095(age+1)+2.645} = e^{0.095age} e^{0.095} e^{2.645}$$

**(after increase) / (original)**

$$\frac{e^{0.095age} e^{2.645} e^{0.095}}{e^{0.095age} e^{2.645}} = e^{0.095} = 1.099$$

When age add 1, the odds ration of have cancer increase 1.099

# Advanced - WOE Encoding (Only for Binary Classification)

Weight of Evidence

$$WoE = \ln\left(\frac{\% non - events}{\% events}\right)$$

To avoid division by zero

$$WoE_{adj} = \ln\left(\frac{\text{Number of non-events in a group} + 0.5 / \text{Number of non-events}}{\text{Number of events in a group} + 0.5 / \text{Number of events}}\right)$$



# Advanced - WOE Encoding (Only for Binary Classification)

## Weight of Evidence

Feature	Outcome	WoE
A	1	0.4
A	0	0.4
A	1	0.4
A	1	0.4
B	1	0.74
B	1	0.74
B	0	0.74
C	1	-0.35
C	1	-0.35

	Non-events	Events	% of non-events	% of events	WoE
A	1	3	50	42	$\ln\left(\frac{(1 + 0.5)/_2}{(3 + 0.5)/_7}\right) = 0.4$
B	1	2	50	29	$\ln\left(\frac{(1 + 0.5)/_2}{(2 + 0.5)/_7}\right) = 0.74$
C	0	2	0	29	$\ln\left(\frac{(0 + 0.5)/_2}{(2 + 0.5)/_7}\right) = -0.35$

# Advanced - WOE Encoding (Only for Binary Classification)

## Function `do_woe_encoding`

```
total_vc = data[label].value_counts().sort_index()
def woe(pipe, total_vc):
    # Count by label in this group
    group_vc = pipe[label].value_counts().sort_index()

    # Some class in the feature is missing, fill zero to missing class
    if len(group_vc) < len(total_vc):
        for key in total_vc.index:
            if key not in group_vc:
                group_vc[key] = 0.
        group_vc = group_vc.sort_index()

    # WOE formula
    r = ((group_vc + 0.5) / total_vc).values

    # Odd ratio => 1 to 0, you can define meaning of each class
    return np.log(r[1] / r[0])

return data.groupby(x).apply(lambda pipe: woe(pipe, total_vc))
```

	Group dist	Global dist
0	0	2
1	2	7

# Advanced - Target Encoding

## Nominal Variable Frequency Encoding

Feature	Encoded Feature
A	0.44
A	0.44
A	0.44
A	0.44
B	0.33
B	0.33
B	0.33
C	0.22
C	0.22

A	0.44 (4 out of 9)
B	0.33 (3 out of 9)
C	0.22 (2 out of 9)

# Advanced - Target Encoding

## Nominal Variable Mean Encoding

Feature	Outcome	MeanEncode
A	1	0.75
A	0	0.75
A	1	0.75
A	1	0.75
B	1	0.66
B	1	0.66
B	0	0.66
C	1	1.00
C	1	1.00

A	0.75 (3 out of 4)
B	0.66 (2 out of 3)
C	1.00 (2 out of 2)

## Topic : WOE + Target Encoding

<b>Filename</b>	lab_eda_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	<ul style="list-style-type: none"><li>→ Add WOE Encoding Feature</li><li>→ Add Target Encoding Feature</li></ul>
<b>Duration</b>	About 15 min


# Advanced - Entropy Encoding

Entropy  $\Rightarrow$  Define the events first!

$$-\sum_i p_i \times \log(p_i)$$

Lower entropy give more information !

Flip Coin Example



# Advanced - Entropy Encoding

## Entropy

Feature	Outcome
A	1
A	0
A	1
A	1
B	1
B	1
B	0
C	1
C	1

Event: 0, 1

$$\text{entropy}(A) = - (1/4 \times \log(1/4) + 3/4 \times \log(3/4)) \\ = 0.81$$

$$\text{proba}(A) = 4/9$$


$$\Rightarrow \text{result} = 0.81 \times 4/9 = 0.36$$

$$\text{entropy}(C) = -(1 \times \log(1)) = 0$$


# Advanced - Polynomial Encoding




$$Y = W_1X + b$$



Normal



$$Y = W_1X + W_2X^2 + W_2X^3 + b$$



Better fitting capability

```
from keras.layers import Dense
```

```
nets = Dense(units=64, ...)  
nets = Dense(units=32, ...)  
logits = Dense(units=1, ...)
```

[Tensorflow Playground](#)




## Topic : Entropy Encoding + Polynomial Encoding

<b>Filename</b>	lab_eda_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	<ul style="list-style-type: none"><li>→ Add Entropy encoding</li><li>→ Add Polynomial encoding</li></ul>
<b>Duration</b>	About 10 min

# Advanced - PCA (Principal Component Analysis)

Sometimes we could face the curse of dimensionality




- Reduce the dimension of data
- Coordinate system transformation
- Mutually orthogonal axis
- Linear transformation

```
from sklearn.decomposition import PCA  
  
pca = PCA(32)  
data = pca.fit_transform(data)
```

So simple ! thank god we have  
scikit learn


# Advanced - AutoEncoder



Minimize Loss(input, output)

- **Goal**
  - Learn a specified data representation
  - Reduce the dimension of data
- Unsupervised learning
  - The input is the label
- Can be non-linear transformation
- The "**Code**" is what we want

# Advanced - AutoEncoder



	<b>AutoEncoder</b>	<b>PCA</b>
<b>Linear?</b>	Non-linear	Linear
<b>Dimension limitation</b>	Non-limited	Less than input dimension

# Advanced - AutoEncoder

```
inputs = Input(shape=(inputs_dim, ))
# Encoder
encoded = Dense(inputs_dim, activation='selu')(inputs)
encoded = Dense(128, activation='selu')(encoded)
encoded = Dense(64, activation='selu')(encoded)

encoded = Dense(64, activation='selu')(encoded)

# Decoder
decoded = Dense(64, activation='selu')(encoded)
decoded = Dense(128, activation='selu')(decoded)
decoded = Dense(inputs_dim, activation='linear')(decoded)

# this model maps an input to its reconstruction
autoencoder = Model(inputs, decoded)
# Adam Optimizer + Mean square error loss
autoencoder.compile(optimizer='adam', loss='mse')

# this model maps an input to its encoded representation
encoder = Model(inputs, encoded)
```

Coded

AutoEncoder model  
for training


Encoder model for  
prediction



## Topic : PCA Encoding + AutoEncoder Encoding

<b>Filename</b>	lab_eda_finance_customer_churn.ipynb
<b>Data</b>	Financial Customer Churn Prediction
<b>Target</b>	<ul style="list-style-type: none"><li>→ Add PCA Encoding</li><li>→ Add AutoEncoder Encoding</li></ul>
<b>Duration</b>	About 10 min

# Advanced - Imbalanced Data For Classification



`class_weight: ...`

*This can be useful to tell the model to "pay more attention" to samples from an under-represented class.*

$$\text{Loss} = - \left( \underbrace{W * Y \log(Y^{\wedge})}_{\text{Minority}} + \underbrace{(1 - Y) \log(1 - Y^{\wedge})}_{\text{Majority}} \right)$$

Minority : Majority = 1 : 40

class weights  $\Rightarrow$  {Minority : 40, Majority: 1}

# Advanced - RFM (Recency, Frequency, Monetary)

How recently, how often and how much did they buy.

Train	Company	rfm_all_freq	distress_num	rfm_all_mean
	36	1	0.026703	0.026703
	36	2	0.020268	0.023485
	36	3	-0.046938	0.000011
	36	4	-0.290000	-0.072492
	36	5	-0.447700	-0.147533
	36	6	-0.333620	-0.178548
Test	Company	rfm_all_freq	distress_num	rfm_all_mean
	36	6	?	-0.178548
	36	6	?	-0.178548
	36	6	?	-0.178548

Beware the "Data leakage", label not in test data, so we take the RFM value from the last moment of train data





# Conclusion

- Domain knowledge is still the key of model performance  
⇒ Why do you know RFM are good for transactional data?
- Deep learning can learning the feature transformation, but still got limitation
- Still need “a little” trial and error